A Wasserstein perspective on inference

Luca Raffo
EPFL, Institute of Mathematics - luca.raffo@epfl.ch
Capital Fund Management - luca.raffo@cfm.com

June 2025

Contents

1 Introduction

Score matching, introduced in [3], yields a new heuristic to estimate continuous statistical models where the probability density function is known only up to a multiplicative normalization constant. The method is shown to be locally consistent under identifiability of the model, and the estimation does

not require to compute the normalization constant.

Here we propose a new point of view that shows that score matching is equivalent to searching for the model that minimizes the Wasserstein gradient (under empirical expectation) of the KL divergence between the real density and the estimated one.

Then we move on to show a deep connection between score matching and minimum probability flow. The latter, introduced in [4], is a different method for estimating statistical models, initially developed for discrete domains such as Ising models, but easily adapted to continuous problems.

This framework can be generalized to different functionals and may lead to new methods for parametric statistical inference. We show and example with MMD divergences.

2 Background on Score Matching

As in usual statistical inference frameworks, we start from a given set of datapoints $D = \{x^{(1)}, ..., x^{(n)}\}$ with $x^{(i)} \in \mathbb{R}^d$, sampled via $X^{(1)}, ..., X^{(n)} \overset{\text{iid}}{\sim} \mu_g$, where μ_g is the real (unknown) distribution whose density with respect to the Lebesgue measure is g.
We want to model the distribution with π_θ, absolutely continuous with respect to Lebesgue, and with density $f_\theta(x) = \frac{1}{Z} e^{-H_\theta(x)}$.

Score matching reads

$$\begin{aligned}\theta_{SM} &= \arg\min_\theta \mathbb{E}_g[\|\nabla_x \log f_\theta - \nabla_x \log g\|_2^2] \\ &= \arg\min_\theta \mathbb{E}_g[(\nabla_x H_\theta)^2 - 2\Delta_x H_\theta] \text{ after integrating by parts},\end{aligned}$$

whose sample version is

$$\begin{aligned}\theta^*_{SM} &= \arg\min_\theta \mathbb{E}_D[\|\nabla_x \log f_\theta - \nabla_x \log g\|_2^2] \\ &= \arg\min_\theta \mathbb{E}_D[(\nabla_x H_\theta)^2 - 2\Delta_x H_\theta],\end{aligned}$$

which does not require the normalization constant Z and can be computed from the data.

In [3] it is shown that θ_{SM} is locally consistent if the model is identifiable (i.e. under the condition that if $\theta_1 \neq \theta_2$ then $\pi_{\theta_1} \neq \pi_{\theta_2}$), and the sample version is asymptotically equivalent to the population one due to the strong law of large numbers.

3 Background on Wasserstein Gradient Flows

The main objective of this section is to unify the notation regarding flows of measures and to define properly Wasserstein gradient flows.

3.1 Flows of measures

Let us sample $X_0 \sim \mu_0$, with $d\mu_0 = f_0\, d\lambda$ and let $v_t : \mathbb{R}^d \to \mathbb{R}$ be any vector field. If we evolve our particle via

$$\dot{X}_t = v_t(X_t), \tag{1}$$

we find out that $\mu_t := Law(X_t)$ is absolutely continuous with respect to Lebesgue, it has finite second moment, and its density satisfies the continuity equation, i.e. it satisfies

$$\partial_t f_t + \nabla \cdot (v_t f_t) = 0 \tag{2}$$

in weak sense. The proof can be found in [1].

3.2 Background on Wasserstein spaces

Consult manuscript I.

3.3 Wasserstein gradient flows

Given a functional $\mathcal{F} : \mathcal{P}_2^{ac}(\lambda) \to \mathbb{R}$ with bounded first variation, we define its Wasserstein gradient at $\mu \in \mathcal{P}_2^{ac}(\lambda)$ as

$$\begin{aligned} \nabla_{\mathcal{W}_2}\mathcal{F}(\mu) : \mathbb{R}^d &\to \mathbb{R}^d \\ x &\mapsto \nabla[\delta\mathcal{F}(\mu)](x). \end{aligned}$$

Now we fix a functional $\mathcal{F}$ with bounded first variation, and use $-\nabla_{\mathcal{W}_2}\mathcal{F}(\mu_t)$ as our vector field v_t in (1), so that μ_t will evolve via

$$\partial_t f_t + \nabla \cdot (-\nabla_{\mathcal{W}_2}\mathcal{F}(\mu_t) f_t) = 0, \tag{3}$$

that is known as *Wasserstein gradient flow* of μ_t with respect to $\mathcal{F}$, started at μ_0.

There are many vector fields v_t such that the ODE (1) induces the PDE (3), it turns out that the *most economical* one, i.e. the one which minimizes $\|v_t\|^2_{L_2(\mu_t)}$ is

$$v_t = -\nabla_{\mathcal{W}_2}\mathcal{F}(\mu_t). \tag{4}$$

Again, we suggest consulting [1] for a proof.

An important property about Wasserstein gradient flows, in similar analogy with gradient flows in finite dimensional spaces, is that if μ_t is following the gradient flow with respect fo $\mathcal{F}$, then

$$\frac{d}{dt}\mathcal{F}(\mu_t) = -\langle \nabla_{\mathcal{W}_2}\mathcal{F}(\mu_t), \nabla_{\mathcal{W}_2}\mathcal{F}(\mu_t)\rangle_{L_2(\mu_t)} \tag{5}$$

$$= -\|\nabla_{\mathcal{W}_2}\mathcal{F}(\mu_t)\|_{L_2(\mu_t)}, \tag{6}$$

showing that these gradient flows dissipate energy along the flow, providing a principled approach to minimizing functionals.

Clearly if a functional $\mathcal{F}$ is strongly displacement convex, then it has a unique minima μ^* (to see this, if there were two such minima, we can arrive to a contradiction by considering the geodesic interpolating between the two of them).

The main result for this section, denoted in [2] as Poljak–Łojasiewicz inequality, states that the Wasserstein gradient flow with respect to a strongly displacement convex $\mathcal{F}$, started at any $\mu_0 \in \mathcal{P}_2^{ac}(\lambda)$, converges exponentially fast towards the unique minimizer $\mu^* \in \mathcal{P}_2^{ac}(\lambda)$.

It turns out that $\mathcal{F}(\cdot) = \mathcal{D}_{KL}(\cdot\|\pi_\theta)$ is (strongly) displacement convex, so it is reasonable trying to minimize it via Wasserstein gradient flows.

4 Score Matching and Wasserstein Gradient Flows

Let us come back to our usual framework of μ_g and π_θ. The idea is that if $\mu_g \approx \pi_\theta$, then

$$\mathcal{D}_{KL}(\mu_g \| \pi_\theta) \approx 0, \tag{7}$$

so that the Wasserstein gradient flow of $\mathcal{D}_{KL}(\cdot\|\pi_\theta)$ (which is playing the role of $\mathcal{F}(\cdot)$) started from $\mu_0 := \mu_g$ will be almost stationary. For a sample $X_0 \sim \mu_g$, since $\dot{X}_0 = -\nabla_{\mathcal{W}_2}\mathcal{D}_{KL}(\cdot||\pi_\theta)|_{\mu_0}(X_0)$ is the starting point of the *most economical* ODE inducing (3), the requirement (7) naturally translates into finding

$$\theta^* := \arg\min_\theta \mathbb{E}_g[\| - \nabla_{\mathcal{W}_2}\mathcal{D}_{KL}(\cdot\|\pi_\theta)|_{\mu_0}\|_2^2]. \tag{8}$$

This can also be derived from (5) and (6), asking for $|\frac{d}{dt}\mathcal{D}_{KL}(\cdot\|\pi_\theta)|_{\mu_0}| \approx 0$.

We know that $\nabla_{\mathcal{W}_2}\mathcal{D}_{KL}(\cdot\|\pi_\theta)|_{\mu_0} = \nabla_x H_\theta + \nabla_x \log g$ (consult manuscript I), so that (8) is aiming to find

$$\begin{aligned}\theta^* &= \arg\min_\theta \mathbb{E}_g[(-\nabla_x H_\theta - \nabla_x \log g)^2] \\ &= \arg\min_\theta \mathbb{E}_g[(\nabla_x \log f_\theta - \nabla_x \log g)^2]\end{aligned}$$

that is precisely the starting point of score matching.

5 Score Matching and Minimum Probability Flow

Minimum probability flow, introduced in [4], is another method for statistical inference. In the continuous framework, it can be summarized as

$$\theta_{MPF} = \arg\min_\theta |\frac{d}{dt}\mathcal{D}_{KL}(\mu_g\|\mu_t)|_{t=0}|,$$

with μ_t following the Fokker-Planck[1] equation with potential π_θ.

It has been shown in [4], using brute force computations, that the function

[1] in practice we use a discretized version.

that minimum probability flow is trying to minimize is the same as score matching; in this section we want to show it from another perspective.

The main idea is that the Wasserstein gradient flow of $\mathcal{D}_{KL}(\cdot\|\pi_\theta)$ is the Fokker–Planck equation with potential H_θ, i.e.

$$\partial_t g_t + \nabla \cdot (-\nabla (\log g_t + H_\theta)\, g_t) = 0 \quad \iff \quad \partial_t g_t = \Delta g_t + \nabla \cdot (g_t \nabla H_\theta),$$

and we can show that score matching is in turn equivalent to minimizing $|\frac{d}{dt}\mathcal{D}_{KL}(\mu_t\|\pi_\theta)|_{t=0}|$, as μ_t satisfies the Fokker Planck equation with potential H_θ. Summarizing briefly, both minimum probability flow and score matching minimize an initial velocity of the KL divergence along the same curve, the difference is that they take the velocity in the second and first entry, respectively.

6 Towards MMD

The reasoning developed so far for score matching and minimum probability flow can be generalized to other statistical discrepancies between probability measures. In particular, we consider the case of the *Maximum Mean Discrepancy* (MMD), a widely used distance in kernel-based inference and generative modeling.

For a detailed introduction to MMD distances and their reproducing-kernel interpretation, we refer to the companion note manuscript II.

Given two measures μ and ν on $\mathbb{R}^d$, and a positive-definite kernel $k : \mathbb{R}^d \times \mathbb{R}^d \to \mathbb{R}$ with reproducing kernel Hilbert space (RKHS) $\mathcal{H}_k$, the squared Maximum Mean Discrepancy is defined as

$$\mathrm{MMD}_k^2(\mu,\nu) := \|\mathbb{E}_{x\sim\mu}[k(x,\cdot)] - \mathbb{E}_{y\sim\nu}[k(y,\cdot)]\|_{\mathcal{H}_k}^2. \tag{9}$$

It can be expanded as

$$\mathrm{MMD}_k^2(\mu,\nu) = \mathbb{E}_{x,x'\sim\mu}[k(x,x')] + \mathbb{E}_{y,y'\sim\nu}[k(y,y')] - 2\,\mathbb{E}_{x\sim\mu,y\sim\nu}[k(x,y)].$$

This quantity vanishes if and only if $\mu = \nu$, provided some regularity conditions regarding k.

Our goal is to define a new inference principle analogous to score matching, by replacing the KL divergence with the MMD functional. Specifically, we want to find the model π_θ such that the Wasserstein gradient of $\mathrm{MMD}_k^2(\cdot, \pi_\theta)$, evaluated at μ_g, is as small as possible.

Our strategy, which involves minimizing the Wasserstein norm of the Wasserstein gradient of the functional, leads to

$$\theta_k^* = \arg\min_\theta \mathbb{E}_{x\sim\mu_g}\left[\left\|\nabla_{\mathcal{W}_2}\mathrm{MMD}_k^2(\cdot,\pi_\theta)\big|_{\mu_g}(x)\right\|_2^2\right]. \tag{10}$$

Intuitively, this corresponds to requiring that μ_g is almost stationary under the Wasserstein gradient flow of the MMD functional towards π_θ, in perfect analogy with what we did for the KL divergence in the score-matching framework.

Our method requires to compute the Wasserstein gradient of MMD_k^2. Taking the first variation with respect to μ, we find

$$\delta\,\mathrm{MMD}_k^2(x) = 2\Big[(k*\mu)(x) - (k*\pi_\theta)(x)\Big], \qquad (k*\nu)(x) := \int_{\mathbb{R}^d} k(x,z)\,d\nu(z).$$

Hence, the Wasserstein gradient is obtained by differentiating with respect to x:

$$\nabla_{\mathcal{W}_2}\mathrm{MMD}_k^2(\mu,\pi_\theta)(x) = \nabla_x\big[\delta\,\mathrm{MMD}_k^2(x)\big] = 2\Big[\nabla_x(k*\mu)(x) - \nabla_x(k*\pi_\theta)(x)\Big]. \tag{11}$$

By exchanging the derivative and the integral, we obtain

$$\nabla_x(k*\nu)(x) = \int_{\mathbb{R}^d} \nabla_x k(x,z)\,d\nu(z),$$

so that

$$\nabla_{\mathcal{W}_2}\mathrm{MMD}_k^2(\mu,\pi_\theta)(x) = 2\left[\int_{\mathbb{R}^d}\nabla_x k(x,z)\,d\mu(z) - \int_{\mathbb{R}^d}\nabla_x k(x,y)\,d\pi_\theta(y)\right]. \tag{12}$$

Evaluating at $\mu = \mu_g$ gives

$$\nabla_{\mathcal{W}_2}\mathrm{MMD}_k^2(\cdot,\pi_\theta)\big|_{\mu_g}(x) = 2\left[\mathbb{E}_{Z\sim\mu_g}[\nabla_x k(x,Z)] - \mathbb{E}_{Y\sim\pi_\theta}[\nabla_x k(x,Y)]\right]. \tag{13}$$

We now investigate explicit forms for common kernels.

1. *Linear kernel:*

$$k(x,y) = x^\top y, \qquad \nabla_x k(x,y) = y.$$

Therefore,

$$\nabla_{\mathcal{W}_2} \mathrm{MMD}_k^2\big|_{\mu_g}(x) = 2\Big[\mathbb{E}_{Z\sim\mu_g}[Z] - \mathbb{E}_{Y\sim\pi_\theta}[Y]\Big],$$

which is constant in x, and the loss reduces to the squared difference between the means:

$$\mathcal{L}_{\mathrm{lin}}(\theta) = 4\left\|\mathbb{E}_{Z\sim\mu_g}[Z] - \mathbb{E}_{Y\sim\pi_\theta}[Y]\right\|_2^2.$$

Hence, the linear kernel version simply enforces mean matching between μ_g and π_θ.

2. *Polynomial kernel:*

$$k(x,y) = (x^\top y + c)^p, \qquad \nabla_x k(x,y) = p\,(x^\top y + c)^{p-1} y.$$

The gradient becomes

$$\nabla_{\mathcal{W}_2} \mathrm{MMD}_k^2\big|_{\mu_g}(x) = 2p\Big[\mathbb{E}_{Z\sim\mu_g}[\,y\,(x^\top Z{+}c)^{p-1}\,] - \mathbb{E}_{Y\sim\pi_\theta}[\,Y\,(x^\top Y{+}c)^{p-1}\,]\Big],$$

and the corresponding loss reads

$$\mathcal{L}_{\mathrm{poly}}(\theta) = 4p^2\,\mathbb{E}_{x\sim\mu_g}\Big\|\mathbb{E}_{Z\sim\mu_g}[\,Z\,(x^\top Z{+}c)^{p-1}\,] - \mathbb{E}_{Y\sim\pi_\theta}[\,Y\,(x^\top Y{+}c)^{p-1}\,]\Big\|_2^2.$$

3. *Gaussian RBF kernel:*

$$k(x,y) = \exp\!\left(-\frac{\|x-y\|^2}{2\sigma^2}\right), \qquad \nabla_x k(x,y) = \frac{y-x}{\sigma^2}\,k(x,y).$$

In this case,

$$\nabla_{\mathcal{W}_2} \mathrm{MMD}_k^2\big|_{\mu_g}(x) = \frac{2}{\sigma^2}\left[\mathbb{E}_{Z\sim\mu_g}\big[(Z-x)\,k(x,Z)\big] - \mathbb{E}_{Y\sim\pi_\theta}\big[(Y-x)\,k(x,Y)\big]\right],$$

and the loss takes the form

$$\mathcal{L}_{\mathrm{RBF}}(\theta) = \frac{4}{\sigma^4}\,\mathbb{E}_{x\sim\mu_g}\Big\|\mathbb{E}_{Z\sim\mu_g}[(Z{-}x)\,k(x,Z)] - \mathbb{E}_{Y\sim\pi_\theta}[(Y{-}x)\,k(x,Y)]\Big\|_2^2.$$

Each choice of kernel therefore yields a different "moment-matching" condition between μ_g and π_θ, ranging from simple mean matching (linear kernel) to high-order nonlinear dependencies (polynomial and Gaussian kernels).

7 Numerical Experiments

We implement our methods to infer the covariance matrix of Gaussian model.

As we saw, minimizing the Wasserstein gradient of the KL divergence led to the score matching loss, while for the MMD we derived three loss functions. The MMD loss with linear kernel will not be informative on the covariance, as it only looks at the mean of the models.

The implementation and results can be found here.

References

[1] L. Ambrosio, N. Gigli, and G. Savare. *Gradient Flows: In Metric Spaces and in the Space of Probability Measures*. Lectures in Mathematics. ETH Zürich. Birkhäuser Basel, 2008.

[2] Sinho Chewi, Jonathan Niles-Weed, and Philippe Rigollet. Statistical optimal transport, 2024.

[3] Aapo Hyvärinen. Estimation of non-normalized statistical models by score matching. *Journal of Machine Learning Research*, 6(24):695–709, 2005.

[4] Jascha Sohl-Dickstein, Peter Battaglino, and Michael Robert DeWeese. A new method for parameter estimation in probabilistic models: Minimum probability flow. *CoRR*, abs/2007.09240, 2020.